

Pawlak_Dominik_HW3

April 13, 2021

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import warnings
import seaborn as sns
warnings.filterwarnings('ignore')
np.random.seed = 44
```

```
[2]: weather_df = pd.read_csv('australia.csv')
weather_df.head()
```

```
[2]:
```

	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustSpeed	\
0	17.9	35.2	0.0	12.0	12.3	48.0	
1	18.4	28.9	0.0	14.8	13.0	37.0	
2	19.4	37.6	0.0	10.8	10.6	46.0	
3	21.9	38.4	0.0	11.4	12.2	31.0	
4	24.2	41.0	0.0	11.2	8.4	35.0	

	WindSpeed9am	WindSpeed3pm	Humidity9am	Humidity3pm	Pressure9am	\
0	6.0	20.0	20.0	13.0	1006.3	
1	19.0	19.0	30.0	8.0	1012.9	
2	30.0	15.0	42.0	22.0	1012.3	
3	6.0	6.0	37.0	22.0	1012.7	
4	17.0	13.0	19.0	15.0	1010.7	

	Pressure3pm	Cloud9am	Cloud3pm	Temp9am	Temp3pm	RainToday	RainTomorrow
0	1004.4	2.0	5.0	26.6	33.4	0	0
1	1012.1	1.0	1.0	20.3	27.0	0	0
2	1009.2	1.0	6.0	28.7	34.9	0	0
3	1009.1	1.0	5.0	29.1	35.6	0	0
4	1007.4	1.0	6.0	33.6	37.6	0	0

```
[3]: X = weather_df.drop(['RainTomorrow'], axis=1)
y = np.array(weather_df['RainTomorrow'])
```

```
[4]: weather_df.describe()
```

```
[4]:
```


	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine \
count	56420.000000	56420.000000	56420.000000	56420.000000	56420.000000
mean	13.464770	24.219206	2.130397	5.503135	7.735626
std	6.416689	6.970676	7.014822	3.696282	3.758153
min	-6.700000	4.100000	0.000000	0.000000	0.000000
25%	8.600000	18.700000	0.000000	2.800000	5.000000
50%	13.200000	23.900000	0.000000	5.000000	8.600000
75%	18.400000	29.700000	0.600000	7.400000	10.700000
max	31.400000	48.100000	206.200000	81.200000	14.500000

	WindGustSpeed	WindSpeed9am	WindSpeed3pm	Humidity9am	Humidity3pm \
count	56420.000000	56420.000000	56420.000000	56420.000000	56420.000000
mean	40.877366	15.667228	19.786778	65.874123	49.601985
std	13.335232	8.317005	8.510180	18.513289	20.197040
min	9.000000	2.000000	2.000000	0.000000	0.000000
25%	31.000000	9.000000	13.000000	55.000000	35.000000
50%	39.000000	15.000000	19.000000	67.000000	50.000000
75%	48.000000	20.000000	26.000000	79.000000	63.000000
max	124.000000	67.000000	76.000000	100.000000	100.000000

	Pressure9am	Pressure3pm	Cloud9am	Cloud3pm	Temp9am \
count	56420.000000	56420.000000	56420.000000	56420.000000	56420.000000
mean	1017.239505	1014.795580	4.241705	4.326515	18.204961
std	6.909357	6.870892	2.797162	2.647251	6.567991
min	980.500000	977.100000	0.000000	0.000000	-0.700000
25%	1012.700000	1010.100000	1.000000	2.000000	13.100000
50%	1017.200000	1014.700000	5.000000	5.000000	17.800000
75%	1021.800000	1019.400000	7.000000	7.000000	23.300000
max	1040.400000	1038.900000	8.000000	9.000000	39.400000

	Temp3pm	RainToday	RainTomorrow
count	56420.000000	56420.000000	56420.000000
mean	22.710333	0.220879	0.220259
std	6.836543	0.414843	0.414425
min	3.700000	0.000000	0.000000
25%	17.400000	0.000000	0.000000
50%	22.400000	0.000000	0.000000
75%	27.900000	0.000000	0.000000
max	46.100000	1.000000	1.000000

```
[5]: weather_df['RainTomorrow'].hist(bins = 3, figsize = (5, 6))
plt.ylabel('Number of cases')
plt.xlabel('Rain Tomorrow')
plt.xticks([0, 1])
plt.show()
```



Zatem klasa “brak deszczu” występuje dużo więcej razy.

```
[6]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳stratify=y)
```

```
[7]: print(X_train.shape, X_test.shape)
```

```
(45136, 17) (11284, 17)
```

1 Logistic Regression

```
[8]: from sklearn.metrics import accuracy_score
from sklearn.model_selection import cross_val_score
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve
```

```

from sklearn.linear_model import LogisticRegression
from sklearn import tree
from sklearn.svm import SVC

modelLR = LogisticRegression(penalty = 'none', random_state=35, max_iter=100000)
modelLR.fit(X_train, y_train)
y_hat = modelLR.predict(X_test)

pred_prob1 = modelLR.predict_proba(X_test)
fpr1, tpr1, _ = roc_curve(y_test, pred_prob1[:,1], pos_label=1)

print(f'Accuracy: {round(accuracy_score(y_test, y_hat), 4)}')
print(f'Cross-Validation mean score: {round(np.mean(cross_val_score(modelLR, X_train, y_train, cv=5)), 4)}')
print(f'ROC-AUC Score: {round(roc_auc_score(y_test, pred_prob1[:,1]), 4)}')

```

Accuracy: 0.8535
Cross-Validation mean score: 0.8526
ROC-AUC Score: 0.8816

2 Decision tree

```

[9]: modelDT = tree.DecisionTreeClassifier(random_state = 35, max_depth = 10)
modelDT.fit(X_train, y_train)
y_hat = modelDT.predict(X_test)

pred_prob2 = modelDT.predict_proba(X_test)
fpr2, tpr2, _ = roc_curve(y_test, pred_prob2[:, 1], pos_label=1)

print(f'Accuracy: {round(accuracy_score(y_test, y_hat), 4)}')
print(f'Cross-Validation mean score: {round(np.mean(cross_val_score(modelDT, X_train, y_train, cv=5)), 4)}')
print(f'ROC-AUC Score: {round(roc_auc_score(y_test, pred_prob2[:, 1]), 4)}')

```

Accuracy: 0.8379
Cross-Validation mean score: 0.8292
ROC-AUC Score: 0.8449

3 SVC

```

[10]: modelSVC = SVC(gamma='scale', degree = 5, kernel = 'poly', max_iter = 100000, random_state = 35)
modelSVC.fit(X_train, y_train)
y_hat = modelSVC.predict(X_test)

fpr3, tpr3, _ = roc_curve(y_test, y_hat, pos_label=1)

```

```

print(f'Accuracy: {round(accuracy_score(y_test, y_hat), 4)}')
print(f'Cross-Validation mean score: {round(np.mean(cross_val_score(modelDT, X_test, y_test)), 4)}')
print(f'ROC-AUC Score: {round(roc_auc_score(y_test, y_hat), 4)}')

```

Accuracy: 0.8513

Cross-Validation mean score: 0.8292

ROC-AUC Score: 0.711


```

[11]: random_probs = [0 for i in range(len(y_test))]
p_fpr, p_tpr, _ = roc_curve(y_test, random_probs, pos_label=1)

plt.plot(fpr1, tpr1, linestyle='--', color='orange', label='Logistic Regression')
plt.plot(fpr2, tpr2, linestyle='--', color='green', label='Decision Tree')
plt.plot(fpr3, tpr3, linestyle='--', color='red', label='SVC')
plt.plot(p_fpr, p_tpr, linestyle='--', color='blue')

plt.title('ROC curve')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive rate')
plt.legend(loc='best')
plt.show();

```



Po wykresie krzywej ROC oraz przedstawionych wynikach innych miar jakości klasyfikatora, widzimy, że z wybranych przeze mnie modeli najlepiej radzi sobie model regresji logistycznej.