

PD4

May 11, 2021

1 PD4 - Jan Smoleń

```
[1]: import pandas as pd
import numpy as np
from matplotlib import pyplot as plt
import seaborn as sns
import dalex as dx
import pickle
np.random.seed = 46
import shap
import xgboost as xgb
from sklearn.metrics import accuracy_score
import warnings
import plotly
warnings.filterwarnings('ignore')
from sklearn.svm import SVC
from sklearn.model_selection import GridSearchCV
```

1.1 Wczytywanie danych

```
[2]: aps=dx.datasets.load_apartments()
from sklearn.preprocessing import OneHotEncoder
aps=dx.datasets.load_apartments()
Xa=aps.drop("district", axis=1)
ya=aps["district"]
aps.head()
```

```
[2]:
```

	m2_price	construction_year	surface	floor	no_rooms	district
1	5897	1953	25	3	1	Srodmiescie
2	1818	1992	143	9	5	Bielany
3	3643	1937	56	1	2	Praga
4	3517	1995	93	7	3	Ochota
5	3013	1992	144	6	5	Mokotow

Ponieważ mamy użyć SVM, to potraktujemy to zadanie jako klasyfikację ze względu na dzielnicę.

```
[3]: wines=pd.read_csv("winequality-red.csv")
wines["is_good"] = wines.apply(lambda row: 1 if row.quality > 5 else 0, axis =
    ↪1)
Xw = wines.drop(["quality", "is_good"], axis = 1)
yw = wines[["is_good"]]
wines.head()
```

```
[3]:      fixed acidity  volatile acidity  citric acid  residual sugar  chlorides \
0           7.4           0.70           0.00           1.9           0.076
1           7.8           0.88           0.00           2.6           0.098
2           7.8           0.76           0.04           2.3           0.092
3          11.2           0.28           0.56           1.9           0.075
4           7.4           0.70           0.00           1.9           0.076

      free sulfur dioxide  total sulfur dioxide  density    pH  sulphates \
0           11.0           34.0    0.9978  3.51           0.56
1           25.0           67.0    0.9968  3.20           0.68
2           15.0           54.0    0.9970  3.26           0.65
3           17.0           60.0    0.9980  3.16           0.58
4           11.0           34.0    0.9978  3.51           0.56

      alcohol  quality  is_good
0         9.4         5         0
1         9.8         5         0
2         9.8         5         0
3         9.8         6         1
4         9.4         5         0
```

Ramkę danych o winie potraktujemy jako zadanie klasyfikacji - czy wino jest dobre (ma ocenę powyżej 5) czy nie.

1.2 Podział na zbiory testowe i treningowe

```
[4]: from sklearn.model_selection import train_test_split
Xw_train, Xw_test, yw_train, yw_test = train_test_split(Xw, yw, test_size = 0.
    ↪2, random_state = 1613)
```

```
[5]: Xa_train, Xa_test, ya_train, ya_test = train_test_split(Xa, ya, test_size = 0.
    ↪2, random_state = 1613)
```

1.3 SVM

```
[6]: from sklearn.svm import SVC
svm_a=SVC()
svm_a.fit(Xa_train, ya_train)
svm_w=SVC()
svm_w.fit(Xw_train, yw_train)
```

[6]: SVC()

1.3.1 Bazowe wyniki, bez standaryzacji

```
[7]: from sklearn.metrics import mean_squared_error  
from sklearn.metrics import accuracy_score
```

```
[8]: ya_preds=svm_a.predict(Xa_test)  
accuracy_score(ya_test, ya_preds)
```

[8]: 0.195

```
[9]: yw_preds=svm_w.predict(Xw_test)  
accuracy_score(yw_test, yw_preds)
```

[9]: 0.6625

Bazowe SVM na surowych zbiorach osiąga bardzo słabe wyniki.

1.3.2 Po standaryzacji

```
[10]: Xa=(Xa-Xa.mean())/Xa.std()  
Xw=(Xw-Xw.mean())/Xw.std()  
Xw_train, Xw_test, yw_train, yw_test = train_test_split(Xw, yw, test_size = 0.  
↪2, random_state = 1613)  
Xa_train, Xa_test, ya_train, ya_test = train_test_split(Xa, ya, test_size = 0.  
↪2, random_state = 1613)  
svm_a.fit(Xa_train, ya_train)  
svm_w.fit(Xw_train, yw_train)
```

[10]: SVC()

```
[11]: ya_preds=svm_a.predict(Xa_test)  
accuracy_score(ya_test, ya_preds)
```

[11]: 0.295

```
[12]: yw_preds=svm_w.predict(Xw_test)  
accuracy_score(yw_test, yw_preds)
```

[12]: 0.76875

Samo standaryzowanie bardzo polepszyło wyniki naszych modeli - o ponad 10%.

1.4 Trening

```
[18]: svm_a_tuned=SVC(random_state=42)
c=[] # wartości parametru C
gamma=[] #wartości parametru gamma
for i in range(-4, 5): # orientacyjne wartości na podstawie informacji
    ↪znalezionych w internecie
    c.append(10**i)
for i in range(-4, 5):
    gamma.append(10**i)
gamma.append("auto")
gamma.append("scale")
params = [{'C': c,
            'gamma': gamma,
            'kernel': ["rbf", "linear"]}]}
from sklearn.model_selection import RandomizedSearchCV
rs_svm_a=RandomizedSearchCV(svm_a_tuned, param_distributions=params,
    ↪scoring='accuracy', cv=4, n_jobs=2)
#gs_svm=GridSearchCV(svm_a_tuned, param_grid=params, scoring='accuracy', cv=4,
    ↪n_jobs=2)
rs_svm_a.fit(Xa_train, ya_train)
rs_svm_a.best_params_
```

```
[18]: {'kernel': 'rbf', 'gamma': 0.01, 'C': 10000}
```

```
[19]: rs_svm_a_acc=accuracy_score(gs_svm.predict(Xa_test),ya_test)
rs_svm_a_acc
```

```
[19]: 0.285
```

```
[23]: svm_w_tuned=SVC(random_state=42)
c=[] # wartości parametru C
gamma=[] #wartości parametru gamma
for i in range(-4, 5): # orientacyjne wartości na podstawie informacji
    ↪znalezionych w internecie
    c.append(10**i)
for i in range(-4, 5):
    gamma.append(10**i)
gamma.append("auto")
gamma.append("scale")
params = [{'C': c,
            'gamma': gamma,
            'kernel': ["rbf", "linear"]}]}
rs_svm_w=RandomizedSearchCV(svm_w_tuned, param_distributions=params,
    ↪scoring='accuracy', cv=4, n_jobs=2)
rs_svm_w.fit(Xw_train, yw_train)
rs_svm_w.best_params_
```

```
[23]: {'kernel': 'rbf', 'gamma': 'scale', 'C': 10}
```

```
[24]: rs_svm_w_acc=accuracy_score(gs_w_svm.predict(Xw_test),yw_test)
      rs_svm_w_acc
```

```
[24]: 0.76875
```

Po dosyć długim czasie trenowania, dostaliśmy wyniki takie same albo nawet nieznacznie gorsze od bazowego SVM na wystandaryzowanych danych.

```
[ ]:
```