

Untitled3

May 11, 2021

```
[1]: import dalex as dx
import numpy as np
import pandas as pd

from sklearn.datasets import load_boston
from sklearn.model_selection import train_test_split, RandomizedSearchCV
from sklearn.svm import SVR
from sklearn.metrics import r2_score, mean_squared_error
from sklearn.preprocessing import MinMaxScaler

import warnings
warnings.filterwarnings('ignore')
```

```
[2]: dalex_df = dx.datasets.load_apartments()
dalex_df.head()
```

```
[2]:   m2_price  construction_year  surface  floor  no_rooms  district
1      5897             1953       25      3         1  Srodmiescie
2      1818             1992      143      9         5    Bielany
3      3643             1937       56      1         2      Praga
4      3517             1995       93      7         3      Ochota
5      3013             1992      144      6         5    Mokotow
```

```
[3]: dalex_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1000 entries, 1 to 1000
Data columns (total 6 columns):
#   Column              Non-Null Count  Dtype
---  -
0   m2_price             1000 non-null   int64
1   construction_year    1000 non-null   int64
2   surface              1000 non-null   int64
3   floor                1000 non-null   int64
4   no_rooms             1000 non-null   int64
5   district             1000 non-null   object
dtypes: int64(5), object(1)
memory usage: 54.7+ KB
```

Jako drugi zbiór danych wziąłem zbiór dotyczący mieszkań w Bostonie z Lab1.

```
[4]: boston_dict = load_boston()
      boston_df = pd.DataFrame(boston_dict['data'],
      ↪columns=boston_dict['feature_names'])
      boston_df['MEDV'] = boston_dict['target']

      boston_df.head()
```

```
[4]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	

	PTRATIO	B	LSTAT	MEDV
0	15.3	396.90	4.98	24.0
1	17.8	396.90	9.14	21.6
2	17.8	392.83	4.03	34.7
3	18.7	394.63	2.94	33.4
4	18.7	396.90	5.33	36.2

```
[5]: boston_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   CRIM        506 non-null    float64
1   ZN          506 non-null    float64
2   INDUS       506 non-null    float64
3   CHAS        506 non-null    float64
4   NOX         506 non-null    float64
5   RM          506 non-null    float64
6   AGE         506 non-null    float64
7   DIS         506 non-null    float64
8   RAD         506 non-null    float64
9   TAX         506 non-null    float64
10  PTRATIO     506 non-null    float64
11  B           506 non-null    float64
12  LSTAT       506 non-null    float64
13  MEDV        506 non-null    float64
dtypes: float64(14)
memory usage: 55.5 KB
```

```
[6]: len(dalex_df['district'].unique())
```

[6]: 10

Ponieważ jest tylko 10 kategorii w ramce DALEX, użyjemy one-hot encodingu.

```
[7]: dalex_df_enc = pd.concat([
    pd.get_dummies(dalex_df.district, prefix='District'),
    dalex_df], axis=1).drop(['district'], axis=1)

# zmieńmy jeszcze kolejność kolumn na bardziej intuicyjną

cols = dalex_df_enc.columns.tolist()
cols = cols[-4:] + cols[:-4]

dalex_df_enc = dalex_df_enc[cols]
dalex_df_enc.head()
```

```
[7]:
```

	construction_year	surface	floor	no_rooms	District_Bemowo	\
1	1953	25	3	1	0	
2	1992	143	9	5	0	
3	1937	56	1	2	0	
4	1995	93	7	3	0	
5	1992	144	6	5	0	

	District_Bielany	District_Mokotow	District_Ochota	District_Praga	\
1	0	0	0	0	
2	1	0	0	0	
3	0	0	0	1	
4	0	0	1	0	
5	0	1	0	0	

	District_Srodmiescie	District_Ursus	District_Ursynow	District_Wola	\
1	1	0	0	0	
2	0	0	0	0	
3	0	0	0	0	
4	0	0	0	0	
5	0	0	0	0	

	District_Zoliborz	m2_price
1	0	5897
2	0	1818
3	0	3643
4	0	3517
5	0	3013

```
[8]: X_dalex = dalex_df_enc.drop('m2_price', axis=1)
Y_dalex = dalex_df_enc.m2_price

X_boston = boston_df.drop(['MEDV'], axis=1)
```

```
Y_boston = boston_df['MEDV']

X_train_dalex, X_test_dalex, y_train_dalex, y_test_dalex = train_test_split(
    X_dalex, Y_dalex, test_size = 0.33, random_state = 34)

X_train_boston, X_test_boston, y_train_boston, y_test_boston = train_test_split(
    X_boston, Y_boston, test_size = 0.33, random_state = 34)
```

1 SVM

```
[9]: svm = SVR()
svm.fit(X_train_dalex, y_train_dalex)
y_hat_dalex = svm.predict(X_test_dalex)
print("Dalex")
print("Wynik R2: " + str(r2_score(y_test_dalex, y_hat_dalex)))
print("Miara RMSE: " + str(mean_squared_error(y_test_dalex, y_hat_dalex,
↪squared = False)))
```

Dalex

Wynik R2: -0.0035647663450799616

Miara RMSE: 934.3010814278865

```
[10]: # przeskalujemy nasze dane i ponownie zbudujemy model
scaler = MinMaxScaler()
dalex_df_enc[['construction_year', 'surface', 'floor', 'no_rooms']] = scaler.
↪fit_transform(dalex_df_enc[['
    'construction_year', 'surface', 'floor', 'no_rooms']]])

X_dalex = dalex_df_enc.drop('m2_price', axis=1)
Y_dalex = dalex_df_enc.m2_price

X_train_dalex, X_test_dalex, y_train_dalex, y_test_dalex = train_test_split(
    X_dalex, Y_dalex, test_size = 0.33, random_state = 34)
```

```
[11]: svm = SVR()
svm.fit(X_train_dalex, y_train_dalex)
y_hat_dalex = svm.predict(X_test_dalex)
print("Dalex po przeskalowaniu")
print("Wynik R2: " + str(r2_score(y_test_dalex, y_hat_dalex)))
print("Miara RMSE: " + str(mean_squared_error(y_test_dalex, y_hat_dalex,
↪squared = False)))
```

Dalex po przeskalowaniu

Wynik R2: 0.040746588319345856

Miara RMSE: 913.441681296815

Widzimy, że po przeskalowaniu wyniki modelu uległy poprawieniu. Ten sam eksperyment przeprowadźmy dla datasetu bostońskiego

```
[12]: svm_boston = SVR()
svm_boston.fit(X_train_boston, y_train_boston)
y_hat_boston = svm_boston.predict(X_test_boston)
print("Boston")
print("Wynik R2: " + str(r2_score(y_test_boston, y_hat_boston)))
print("Miara RMSE: " + str(mean_squared_error(y_test_boston, y_hat_boston,
↪squared = False)))
```

Boston

Wynik R2: 0.25006369536003814

Miara RMSE: 7.915412693509835

```
[13]: scaler = MinMaxScaler()
boston_df[['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD',
↪ 'TAX', 'PTRATIO', 'B', 'LSTAT']] = scaler.fit_transform(boston_df[['CRIM',
↪ 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX', 'PTRATIO',
↪ 'B', 'LSTAT']])

X_boston = boston_df.drop('MEDV', axis=1)
Y_boston = boston_df.MEDV

X_train_boston, X_test_boston, y_train_boston, y_test_boston = train_test_split(
    X_boston, Y_boston, test_size = 0.33, random_state = 34)

svm = SVR()
svm.fit(X_train_boston, y_train_boston)
y_hat_boston = svm.predict(X_test_boston)
print("Boston po przeskalowaniu")
print("Wynik R2: " + str(r2_score(y_test_boston, y_hat_boston)))
print("Miara RMSE: " + str(mean_squared_error(y_test_boston, y_hat_boston,
↪squared = False)))
```

Boston po przeskalowaniu

Wynik R2: 0.6089303140466609

Miara RMSE: 5.71595038091061

Wniosek: Skalowanie danych przynosi dobre efekty.

2 Random Search

```
[14]: parameters = dict(
    C = np.arange(start = 0.1, stop = 10000, step = 0.05),
    gamma = ['scale', 'auto'],
    degree = np.arange(1, 80, 1))

svm_rand_dalex = RandomizedSearchCV(svm_boston, parameters, cv=3, n_iter=200)

svm_rand_dalex.fit(X_train_dalex, y_train_dalex)
```

```

print("Najlepsze parametry: " + str(svm_rand_dalex.best_params_))

best_estimator = svm_rand_dalex.best_estimator_
print("Wynik R2: " + str(r2_score(y_test_dalex, best_estimator.
    ↳predict(X_test_dalex))))
print(f'RMSE: {mean_squared_error(y_test_dalex, best_estimator.
    ↳predict(X_test_dalex), squared=False)}')

```

Najlepsze parametry: {'gamma': 'scale', 'degree': 5, 'C': 4608.5500000000002}
 Wynik R2: 0.9708321322696781
 RMSE: 159.28192280776173

```

[15]: svm_rand_boston = RandomizedSearchCV(svm_boston, parameters, cv=3, n_iter=200)

svm_rand_boston.fit(X_train_boston, y_train_boston)
print("Najlepsze parametry: " + str(svm_rand_boston.best_params_))

best_estimator = svm_rand_boston.best_estimator_
print("Wynik R2: " + str(r2_score(y_test_boston, best_estimator.
    ↳predict(X_test_boston))))
print(f'RMSE: {mean_squared_error(y_test_boston, best_estimator.
    ↳predict(X_test_boston), squared=False)}')

```

Najlepsze parametry: {'gamma': 'scale', 'degree': 72, 'C': 196.25000000000006}
 Wynik R2: 0.9054813042861667
 RMSE: 2.810090028250059

[]: