

PD 7

Artur Żółkowski

In [38]:

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import seaborn as sns
from sklearn.metrics import f1_score, precision_score, recall_score
from sklearn.mixture import GaussianMixture
```

In [2]:

```
train_df = pd.read_csv("train.csv")
val_df = pd.read_csv("val.csv")
test_df = pd.read_csv("test.csv")
```

In [4]:

```
print(train_df.shape)
print(val_df.shape)
print(test_df.shape)
```

```
(95, 13)
(17, 13)
(17, 14)
```

In [5]:

```
X_test_df = test_df.iloc[:, 1:]
y_test = test_df.iloc[:, 0]
```

In [6]:

```
train_df
```

Out[6]:

	Alcohol	Malic acid	Ash	Alcalinity of ash	Magnesium	Total phenols	Flavanoids	Nonflavanoid phenols	Proanthoc
0	12.72	1.75	2.28	22.5	84	1.38	1.76	0.48	
1	13.23	3.30	2.28	18.5	98	1.80	0.83	0.61	
2	12.58	1.29	2.10	20.0	103	1.48	0.58	0.53	
3	12.37	1.17	1.92	19.6	78	2.11	2.00	0.27	
4	13.84	4.12	2.38	19.5	89	1.80	0.83	0.48	
...	
90	12.84	2.96	2.61	24.0	101	2.32	0.60	0.53	
91	12.70	3.87	2.40	23.0	101	2.83	2.55	0.43	
92	12.77	2.39	2.28	19.5	86	1.39	0.51	0.48	
93	12.88	2.99	2.40	20.0	104	1.30	1.22	0.24	
94	12.25	4.72	2.54	21.0	89	1.38	0.47	0.53	

95 rows × 13 columns



In [11]:

```
df_all = pd.concat([train_df, val_df, X_test_df])
```

In [12]:

```
df_all.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 129 entries, 0 to 16
```

```
Data columns (total 13 columns):
```

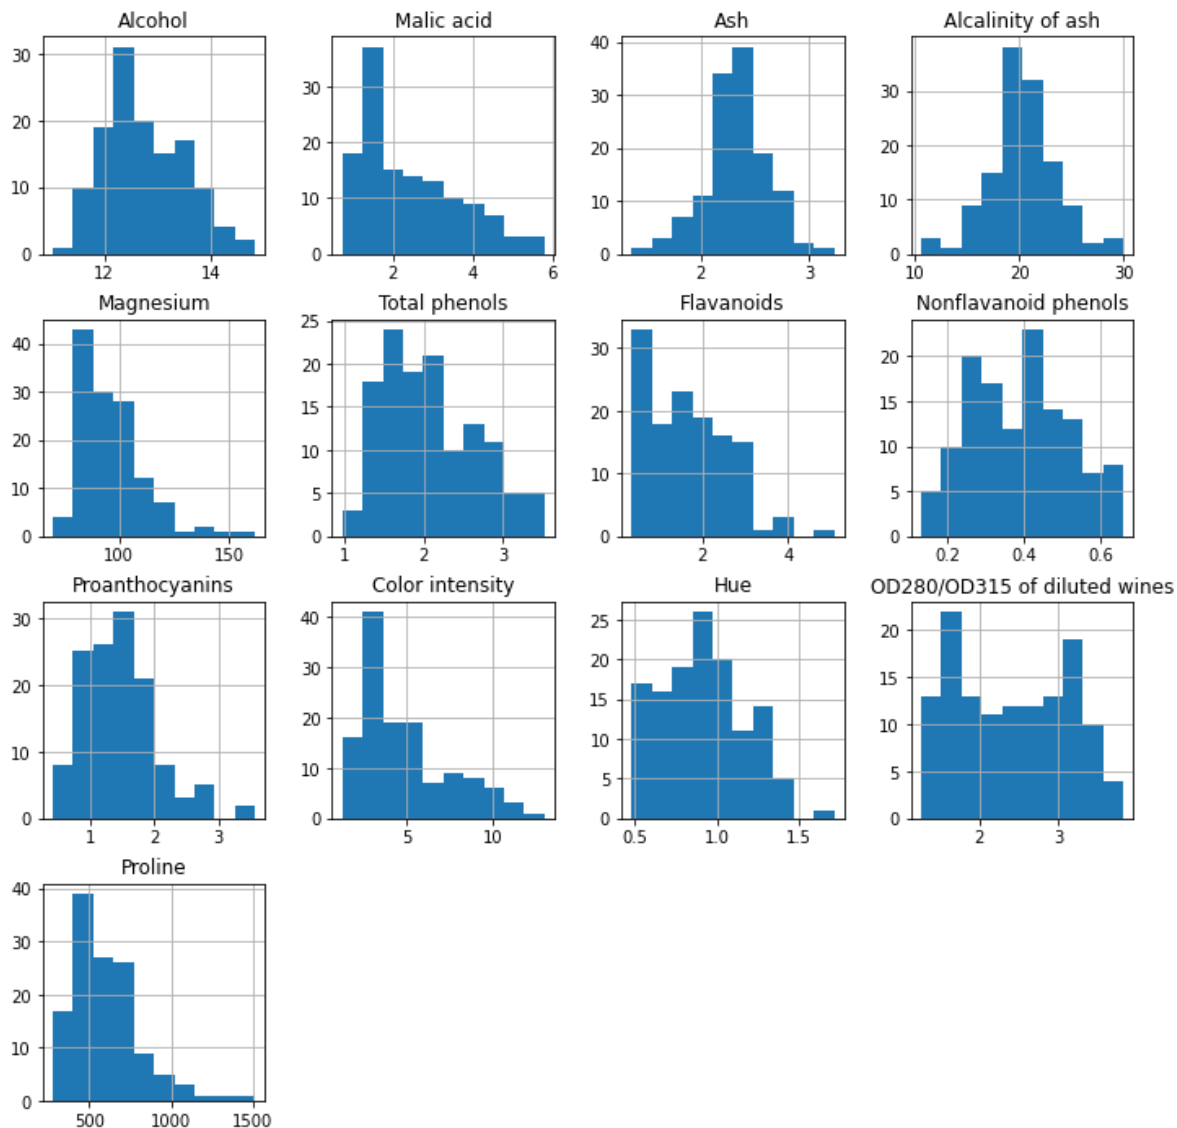
#	Column	Non-Null Count	Dtype
0	Alcohol	129 non-null	float64
1	Malic acid	129 non-null	float64
2	Ash	129 non-null	float64
3	Alcalinity of ash	129 non-null	float64
4	Magnesium	129 non-null	int64
5	Total phenols	129 non-null	float64
6	Flavanoids	129 non-null	float64
7	Nonflavanoid phenols	129 non-null	float64
8	Proanthocyanins	129 non-null	float64
9	Color intensity	129 non-null	float64
10	Hue	129 non-null	float64
11	OD280/OD315 of diluted wines	129 non-null	float64
12	Proline	129 non-null	int64

```
dtypes: float64(11), int64(2)
```

```
memory usage: 14.1 KB
```

In [14]:

```
df_all.hist(figsize=(12,12))  
plt.show()
```




Wykorzystanie GMM do wykrywania próbek odstających

In [15]:

```
bic = []  
for i in range(1, 12):  
    gmm_tmp = GaussianMixture(n_components = i, random_state = 42, covariance_type = "full")  
    gmm_tmp.fit(train_df)  
    bic.append(gmm_tmp.bic(train_df))
```

In [18]:

```
plt.figure(figsize=(14, 8))
plt.plot(range(1, 12), bic)
plt.xlabel("Component number")
plt.ylabel("Bic metric value")
plt.show()
```



In [19]:

```
gmm = GaussianMixture(covariance_type='full', n_components=1, random_state=42)
gmm.fit(train_df)
```

Out[19]:

```
GaussianMixture(random_state=42)
```

Wybór próbek odstających

In [23]:


```
probs_train = gmm.score_samples(train_df)
probs_val = gmm.score_samples(val_df)
```

In [24]:

```
pd.DataFrame(probs_train).hist()
```

Out[24]:

```
array([[<AxesSubplot:title={'center':'0'}>]], dtype=object)
```




In [25]:

```
pd.DataFrame(probs_val).hist()
```

Out[25]:

```
array([[<AxesSubplot:title={'center':'0'}>]], dtype=object)
```



In [28]:

```
probs = pd.concat([pd.DataFrame(zip(probs_val, ["val"]*len(probs_val))), pd.DataFrame(zip(p
probs.reset_index(inplace=True, drop=True)
probs.columns = ["prob", "set"]
```

In [29]:

```
probs
```


Out[29]:

	prob	set
0	-24.109594	val
1	-13.656306	val
2	-18.348802	val
3	-24.731812	val
4	-26.572217	val
...
107	-17.185796	train
108	-16.504880	train
109	-17.588148	train
110	-15.767297	train
111	-18.207360	train

112 rows × 2 columns

In [35]:

```
plt.figure(figsize=(8, 12))
sns.violinplot(data = probs, y = "prob", x = "set", inner = "quartiles")
sns.swarmplot(data = probs, y = "prob", x = "set", color = "black")
plt.show()
```



Bazując na wykresach wybieramy threshold na poziomie -31.

Sprawdzenie wyników

In [36]:

```
test_probs = gmm.score_samples(X_test_df)
y_testh = test_probs < -31
y_testh = y_testh.astype(int)
```

In [39]:

```
results = pd.DataFrame(
    [
        ("precision", precision_score(y_test, y_testh)),
        ("recall", recall_score(y_test, y_testh)),
        ("F1-score", f1_score(y_test, y_testh))
    ],
    columns=["metric", "value"]
)
results
```

Out[39]:

	metric	value
0	precision	0.833333
1	recall	1.000000
2	F1-score	0.909091